

Simple neural networks for text

Quiz, 3 questions

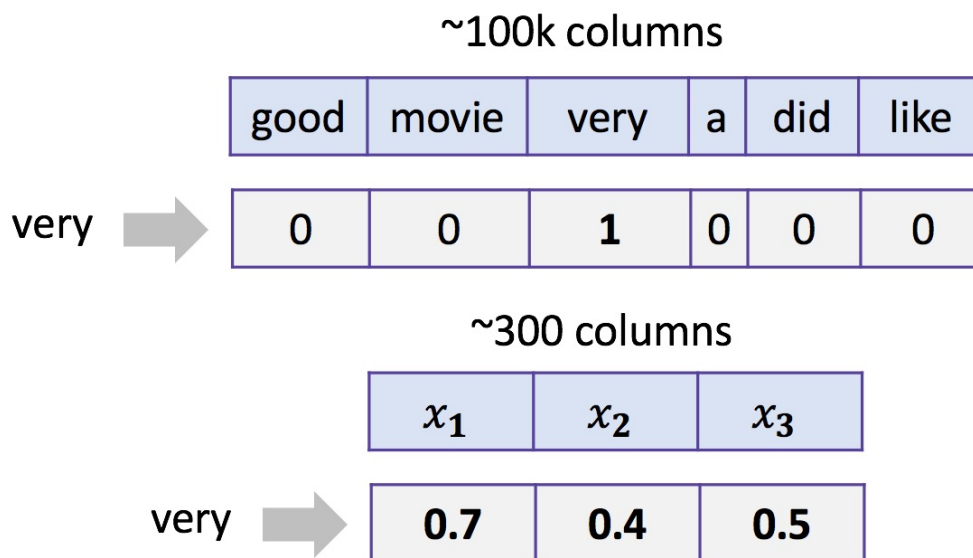
✓ **Congratulations! You passed!**

Next Item



2 / 2
points

1.
Let's recall how we treated words as one-hot sparse vectors in BOW and dense embeddings in neural networks:



Choose correct statements below.



Linear model on top of a **sum** of neural representations can work faster than on top of BOW.



Correct

This is true! We only need to train 300 parameters here. Don't forget to normalize these features row-wise!



You can replace **word2vec** embeddings with any **random** vectors to get a good features descriptor as a **sum** of vectors corresponding to all text tokens.



Un-selected is correct



For **both** word representations we can take a **sum** of vectors corresponding to tokens of any text to obtain good features for this text for further usage in linear model.

Simple neural networks for text

Quiz, 2 questions
Correct
Yes, this is true. Don't forget to normalize these features row-wise!



For **both** word representations we can take a **weighted sum** of vectors corresponding to tokens of any text to obtain good features for this text for further usage in linear model. The **weight** for any token can be an IDF value for that token.

Correct

Yes, this is true. For BOW we effectively get bag of TF-IDF values, where TF is a binary variable. Don't forget to normalize these features row-wise!



2 / 2
points

2.
Let's recall 1D convolutions for words:

Word embeddings

cat	0.7	0.4
sitting	0.2	-0.1
there	-0.5	0.4
or	-0.1	0.8
here	-0.5	0.3

What is the result of 1D convolution + maximum pooling over time for the following kernel **without padding**?

1	0
0	1

0.6

Simple neural networks for text

For Her Response
That's it!
Quiz, 3 questions



1 / 1
point

3.

Let's recall 1D convolutions for characters. Choose correct statements.



One 1D convolutional layer for spotting character 3-grams is enough for solving a practical task.



Un-selected is correct



1D convolutions work better than BOW for huge datasets.



Correct
This is true.



1D convolutions for characters consume one-hot encoded vectors for characters.



Correct
That's right, they are not that long, so this is okay.

